Using Timetags for Program Dependency Enforcement

A. Khalafi, D. Morano, D.R. Kaeli
Northeastern University
dmorano, kaeli@ece.neu.edu

A.K. Uht
University of Rhode Island
uht@ele.uri.edu

Abstract

We discuss how *time tags* can be used for the enforcement of program dependencies. Time tags can serve as the basic ordering enforcement mechanism when large numbers of instructions are executing concurrently. Proposed and future microarchitectures can have hundreds or several hundreds of instructions in flight simultaneously and using standard reservation tags, physical register addresses, and reorder buffers do not scale well for very large instruction windows. The design, use and management of time tags will be discussed. We also provide simulation data for an example microarchitecture that takes advantage of time tags for its dependency ordering.

1 Introduction

A number of studies into the limits of instruction level parallelism (ILP) have been promising in that they have shown that there is parallelism within typical programs [7, 12, 25]. Unfortunately, most of this fine-grained parallelism spans several basic blocks and the relatively small instruction windows (less than about two hundred) of existing processor designs cannot span the program instruction space necessary to begin to exploit this parallelism. A large number of instructions need to be fetched each cycle and executed concurrently in order to expose this ILP. This generally requires a large microarchitecture with a large instruction window and reorder buffer or register undate unit. A fundamental challenge for new microarchitectures will be how to find the inherent instruction level parallelism in programs and then allow execution to occur out of order while still maintaining the architectural program order.

In this paper we discuss a means by which time tags can be used to enforce proper program dependencies while allowing massive out-of-order execution by a suitable (and likely large) microarchitecture. We are proposing to use time tags to maintain and enforce correct program order for all flow dependencies whether they be registers, memory values, or instruction control-flow predicates. It may be noted that the *time* aspect of time tags refers to program-order time and not real time. The goal is to allow for massive speculative out-of-order instruction execution in real time while providing a means of tracking program order time.

Section 2 will provide some high level applications for and background on the use of time tags. Specifically, microarchitectural capabilities that are likely desirable in future machines, and which may be implemented more easily using time tags, are presented. Section 3 presents a description of time tags and how they would fit into a potential microarchitecture. Also discussed are some optional microarchitectural features that can be facilitated once time tags are used as a program dependency enforcement mechanism. Section 4 briefly presents a proposed microarchitecture that uses time tags for all of its dependency tracking and enforcement. Simulated results from this microarchitecture are also presented. Finally, in Section 5 we summarize the current contributions.

2 Application and Background

Current microarchitectures generally employ mechanisms such as a Register Update Unit or Reorder Buffer(s) [4, 21, 10] to enforce proper program dependencies for registers and predicates. Memory dependencies are often

1

enforced through the use of snooping the memory store queue or some other memory disambiguating mechanism [5]. The use of time tags to coordinate and enforce correct program order can be used to eliminate or reduce the need for the severe contention present on either register scoreboards [23] or the architected register file associated with the use of reservation stations [1].

In those microarchitectures that perform speculative execution, there is also the need to access the reorder buffer, which becomes quite problematic as the number of instructions being speculatively executed concurrently grows [17]. Whether speculative instruction operand results are stored in data registers within the reorder buffer or if the results are stored in extra physical registers that hold both architected and temporary values, the contention for the centralized resource is the same. Further, future microarchitectures may employ value speculation [13, 14, 9] as well as control-flow speculation. Time tags would appear to be a good candidate for unifying, to some extent, the tracking of value dependencies along with control-flow dependence.

Several microarchitectures have also been proposed that perform multipath execution. Some of these are described by Wang [28], Uht and Sindagi [25], Heil and Smith [8], the PolyPath by Klauser et al [11], and a simultaneous multithreading oriented approach by Wallace et al [27]. In order to accommodate multipath execution, we employ a *path ID* similar to that described by Skadron and Ahuja [20].

Some newly proposed microarchitectures are also addressing the desire to execute control-independent instructions beyond the join of conditional branches. Some of these are described by Sankaranarayana and Skadron [18, 19], Cher and Vijaykumar [2], and Uht et al [26]. In our present work, we also desire to provide for the speculative execution and reuse of control-flow independent instructions after conditional branches. We will use time tags to facilitate the ordering of operands and instruction execution in this sort of circumstance also.

The Warp Engine [3] used time tags. However their implementation was somewhat cumbersome ; utilizing floating point numbers and machine-wide parameter updating. We will employ a simpler representation of program-order time for our time tags.

3 Time Tags

A time tag is a small integer that uniquely identifies the position of an instruction or an operand in program ordered time. Both instructions and operands are tagged with these values. Typically, time tags take on small positive values to indicate the position of an instruction in program ordered time with respect to the most recently committed and retired instruction, which can be thought of having a time tag value of minus one. The oldest dispatched instruction in flight that is neither yet committed nor squashed would usually have the time tag value of zero. Those instructions in flight, and therefor being in future program ordered time, take on increasingly larger values the further ahead they are in the program ordered instruction stream. Likewise, certain operands that are created by instructions (the outputs of that instruction) would take on the same time value as its instruction. This thus provides the fundamental means by which proper operand dependency order can be tracked and enforced. As instructions, or groups of instructions, having the lowest valued time tags are retired, all of the time tags for all instructions and operands are decremented by the number of instructions just retired. This will keep the next instruction in program-ordered time that is ready to retire next to always have the time tag value of zero. Other variations for the assignment and management of the values are also possible (and have been used) but thinking of the oldest dispatched and not yet retired instruction as having the value zero is most illustrative for this presentation.

In the arrangement described so far, negative valued time tags may still be used depending on the particulars of the microarchitecture that they might be employed for. In such cases, instructions themselves would never take on negative values since there is no need to retain them after they are finally retired. However, operands might still take on negative valued time tags and this would indicate that they have already been committed but that their program-order relationship to each other is desired to be maintained for some purpose. An example might be when memory operands are committed but still need to be snooped in a store queue of some sort. Rather then forcing a FIFO storage arrangement for those committed memory values, it may be more expedient to use time tags for the same purpose.

3.1 Operands

Operands for instructions can be of varying types and different microarchitectures may want to treat each type somewhat differently. At least three types of instruction operands may be identified for most microarchitectures. These are register, memory, and predicate operands. Some microarchitectures may want to allow for more speculative flexibility of certain operands while not allowing as much speculation for other types.

For most microarchitectures, registers are the primary instruction operand and will thus, rightly, deserve the most attention. These register operands would consist of both those that are input to the instruction as well as those that are outputs from the instruction. Typically, all machines with register oriented instruction set architectures (ISAs) could likely benefit from having register operands time tagged.

However, it may be quite advantageous to additionally tag memory and possibly predicate operands as well. The time tagging of memory operands allows them to be treated as flexibly and as speculatively as register operands might be. Memory operands only apply to those instructions that perform memory load/store type functions. A load-type instruction can be thought of having one or more input memory operands and a store-type instruction likewise has one or more output memory operands. With this understanding, load/store instructions can now be handled in ways very similar to all register-only oriented instructions.

For predicate operands, two types need to be distinguished. One type is a part of the ISA of the machine and is therefor visible to programmers. A predicate operand of this type would be used in an ISA such as the iA-64, for example. For our purposes, this sort of explicit predicate operand is identical to a register operand (discussed above) and is simply treated as such. Not as clear is the use of predicate operands that are not a part of the ISA and are therefor not visible to the programmer. This type of predicate in entirely maintained by the microarchitecture itself but still essentially forms an additional input operand for each instruction. This single bit operand is what is used to predicate the instruction's committed execution, the same as if it was explicit in the ISA. This input predicate thus enables or disables its associated instruction from producing its own new output operand. It should be noted that, at any time, any instruction can always still be allowed to execute but that only the output results of that instruction execution must be made available or not depending on the instruction's input predicate operand. For microarchitectures that support these microarchitecture-only predicate operands, they too can be time tagged, thus allowing them the same degree of out-of-order flexibility similar to register or memory operands. Finally, note that even ISAs that sport predicate registers can also independently employ predication of instructions within the microarchitecture.

3.2 Multipath Execution

Not all microarchitectures may uses multiple speculative execution paths but we plan for that case anyway. Some extra complications may arise for those microarchitectures that allow for multiple simultaneous speculative execution paths. There are several ways in which multiple execution paths can be tracked but some care needs to be applied in the context of also using time tags. Generally, a new speculative path may be formed after the encounter of each additional conditional branch in the instructions stream. Execution is not speculative for those instructions just following the last committed instruction but all instructions are speculative after the first unresolved conditional branch is encountered. At least two options are available for assigning time tags to the first instruction following each output of a conditional branch.

One option is to dynamically follow each output path and assign successively higher time tag values to succeeding instructions. Another option is to try to determine if the *taken* outcome of the branch joins with the *not-taken* output instruction stream. If a join is determined, the instruction following the *not-taken* output path can be assigned a time value one higher than the branch itself, while the first instruction on the *taken* path would be assigned whatever value it would have gotten counting up from the first instruction on the *not-taken* path. Further, both options may be employed simultaneously in the same microarchitecture.

In both cases, there will exist instructions in flight that both have the same value for their time tag. Likewise, operands resulting from these instructions would also share the same time-tag value. For our purposes, this ambiguity is resolved through the additional of a *path-ID* value. This path ID would also be tagged to all instructions

and operands just as the time tags are. With the addition of the path-ID, in addition to time tags, a fully unique address space is now possible for all instructions that may be in flight.

3.3   Names and Renaming

We are now in a position to define a name space that uniquely specifies all instructions and operands that may be in flight at any time. For instructions, names for them can be uniquely created with the concatenation of the following components :

- a path ID
- the time tag assigned to a dispatched instruction

For all operands, unique names would consist of :

- type of operand
- a path ID
- time tag
- identifying address

Generally, the type of the operand would be *register*, *memory*, or *predicate*. Remember that ISA visible predicate are treated just as regular data registers within the microarchitecture itself. The use of this component for operand names may not be necessary depending on how the different operand types are transfered among instructions. This is discussed more later. Again, the path ID and time tag are as discussed previously. The identifying address of the operand would differ depending on the type of the operand. For register operands, the identifying address would be the name of the architected register. All ISA architected registers are typically provided a unique numerical address. These would include the general purpose registers, any status or other non-general purpose registers, and any possible ISA predicate registers. For memory operands, the identifying address is just the programmer visible architected memory address of the corresponding memory value. Finally, for predicate operands, the identifying address may be absent entirely, for some microarchitectures, or might be some address that is used within the microarchitecture to further identify the particular predicate register in question. We have explored microarchitecture predication schemes that have used both options above.

Now that all instructions and operands can have unique names within the microarchitecture of a machine, additional options open up for new microarchitectures that were not as obvious or exploitable in previous machines. For example, if each instruction in an issue slot retains a copy of its input or output operands (whatever they may be), effectively a full renaming of all operands will have occurred for all instructions in flight in the machine. All false dependencies are now automatically avoided. There would be no need to limit instruction dispatch or to limit speculative instruction execution due to a limit on the number of non-architected registers available for holding temporary results. Since every possible operand has a unique name and is available to be read in some way, both all needed ordering information and the associated necessary operand values to match a particular ordering would be available. This feature can be used powerfully to eliminate the need for physical register renaming, register update units, or reorder buffers. The sum of the issue slots in the machine would effectively serve as a sort of reorder buffer in themselves. We will next explore how operands are transferred from instruction to instruction while enforcing the necessary true program dependencies for proper execution.

3.4   Operand Forwarding and Snooping

As with microarchitectures not using time tags, operands resulting from the execution of instructions are broadcast forward for use by waiting instructions. As expected, when operands are forwarded, they are done so by transferring not only the identifying address of the operand and its value but also the time tag and a potential path-ID (for those microarchitectures using multipath execution). This tag will be used by subsequent instructions (later in program order time) already dispatched to determine if the operand should be *snarfed* [1] as an input that will trigger the execution of its loaded instruction.

---

[1] snarfing entails snooping address/data buses, and when the desired address value is detected, the associated data value is read
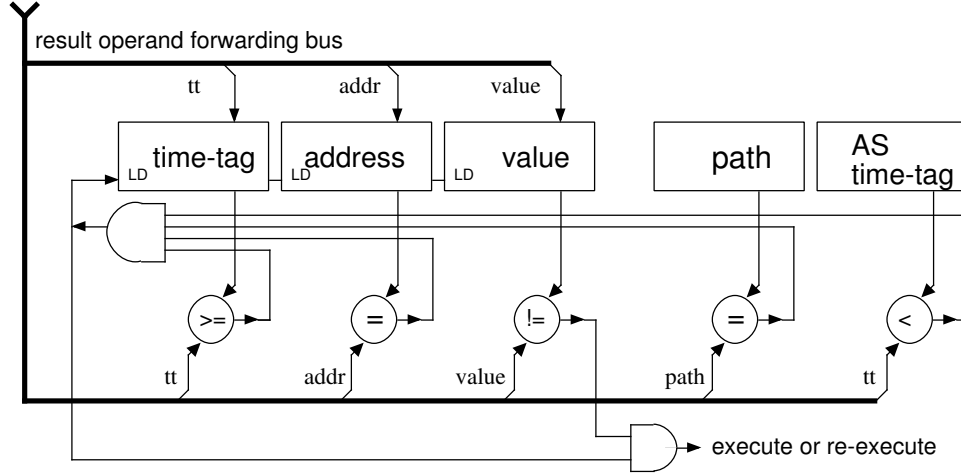
Figure 1: *Instruction Source Operand.* The registers and snooping operation of one of several possible source operands is shown. Just one operand forwarding bus is shown being snooped but typically several operand forwarding buses are snooped simultaneously.

The information associated with each operand that is broadcast from one instruction to subsequent ones is referred to as a *transaction*, and generally consists of :

- transaction type
- operand type
- path ID
- time tag of the originating instruction
- identifying address
- data value for this operand

This above information is typical of all operand transactions. True flow dependencies are enforced through the continuous snooping of these transactions by each dispatched instruction residing in an issue slot that receives the transaction. Each instruction will snoop all operands that are broadcast to it but an operand forwarding interconnect fabric may be devised so that transactions are only sent to those instructions that primarily lie in future program-ordered time from the originating instruction. More information on operand forwarding interconnection networks is presented in a later section. More details on transactions for register, memory, and predicate forward operations are also presented in a later section.

Figure 1 shows the registers inside an issue slot used for the snooping and snarfing of one of its input operands. The *time-tag, address,* and *value* registers are reloaded with new values on each snarf, while the *path* and *instr. time-tag* are only loaded when an instruction is dispatched to the active station. The operand shown is typical for source registers, a source memory operand, or an instruction execution predicate register.

If the path ID and the identifying address of the operand matches any of its current input operands, the instruction then checks if the time tag value is less than its own assigned time tag, and greater than or equal to the time tag value of the last operand that it snarfed, if any. If the snooped data value is different than the input operand data value that the active station already has, a re-execution of the instruction is initiated. This simple rule will enforce all proper flow dependencies while allowing for the potential for massive concurrency to occur.

3.5   Result Forwarding Buses and Operand Filtering

The basic requirement of the interconnection fabric is that it must be able to transport operand results from any instruction to those instructions with higher valued time tags. This corresponds to the forwarding of operands into future program ordered time. There are several choices for a suitable interconnection fabric between instruction

issues slots. A simple interconnection fabric could consist of one or more buses that simply interconnect all issue slots. Although appropriate for smaller sized microarchitectures, this arrangement does not scale as well as some other alternatives. A more appropriate interconnection fabric that would allow for the physical scaling of the microarchitecture may be one in which forwarding buses are segmented with active repeaters between the stages. This arrangement exploits the fact that, for example, register lifetimes are limited [6, 22]. Registers being forwarded to instructions lying close in future program-ordered time will get their operands quickly while those instructions lying beyond the repeater units will incur additional delays. In addition to allowing for physical scaling, it also offers the opportunity for filtering out some operands that do not need to be forwarded beyond a certain point. Another possibility for the operand forwarding fabric is to have separate buses for each type of operand. This sort of arrangement could be used to tailor the available bandwidth provided for each type of operand. It could also allow for a different interconnection network to be used for each type of operand also. We have explored several of these possible variations already.

The opportunity to provide active repeaters between forwarding bus segments also opens up a range of new microarchitectural ideas not easily accomplished without the use of time tagging. Different types of active repeaters can be optionally used. Further, different types of repeaters can be used for the different types of operands. Some can just provide a store-and-forward function while another variation could also store recently forwarded operands and filter out (not forward) newly snooped operands that have already been forwarded previously. The latter type of forwarding repeater unit is termed a *filter unit*. This feature can be used to reduce the operand traffic on the forwarding fabric and thus reduce implementation costs. For example, for memory operands, a cache (call it a *L0 cache*) can be situated inside a *memory filter unit* (MFU) and can hold recently requested memory operands. This can reduce the demand on the L1 data cache and the rest of the memory hierarchy by satisfying some percent of memory operand loads from these caches alone. Register filter units (RFUs) are also possible and can reduce register transfer traffic similarly to the MFUs.

3.6   Operand Forwarding Strategies and Bus Transactions

Although we have so far described the operand forwarding mechanism in simple terms as being the broadcasting of it to those ASes with higher valued time tags, there are some additional details that need to be addressed for a correctly working forwarding solution. These details also differ depending on whether operands for register, memory, or predicates need to be forwarded. There are many possible strategies for forwarding of operands (and of operands of differing types). We now briefly outline three such strategies. One of these is suitable for registers. Another is suitable for registers or memory operands. The third is oriented for the forwarding of predicates. These three strategies are termed *relay forwarding*, *nullify forwarding*, and *predicate forwarding* respectively. In general, each forwarding strategy employs bus transactions of one or more types to implement its complete forwarding solution. The particulars of these transactions are described for each of the forwarding strategies.

3.6.1   Relay Forwarding

This forwarding strategy is quite simple but is also entirely adequate for the forwarding of register operands. In this strategy, when a new register operand needs to be forwarded from an active station, the standard operand information, as previously described in a general, is packaged up into what is termed a *register store* transaction. This transaction type consists of :
- a transaction ID of *register store*
- the path ID on which this instruction is executing
- the time tag of the originating active station
- the register address
- the value of the register

A request is made to arbitrate for an outgoing forwarding bus and this transaction is placed on the bus when it becomes available.

When the instruction associated with an active station gets a new input operand, it will re-execute producing a new output operand. In this forwarding strategy, the new output operand is both stored locally within the

active station itself and is sent out on the outgoing forwarding buses to subsequent (higher time-tag valued) active stations. Previous values of the instruction's output operand is also snooped as if it was an input operand and is also stored locally within the active station. It should also be noted that if the enabling execution predicate for the current instruction changes, either from being enabled to disabled or visa versa, a new output operand is forwarded. If the instruction predicate changes from disabled to enabled, the output operand that was computed by the instruction os forwarded. If the instruction predicate changes from enabled to disabled, the previous value of the output operand (before being changed due to the instruction execution) is forwarded. That previous value is available to the active station because it gets snooped as if it was an additional input. Newly forwarded operands will always superseded and previously forwarded operands. With this strategy, instructions that are located in the program ordered future will eventually always get the correct value that will end up being the committed value if the current instruction ends up being committed itself (ends up being predicated to execute). This is an elegant forwarding strategy and the simplest of the forwarding strategies investigated so far, and is a reasonable choice for the handling register operands. The inclusion of the time tag in the transaction is the key element that allows for the correct ordering of dependencies in the committed program.

3.6.2 Nullify Forwarding

There are limitations to the applicability of the previously discussed forwarding strategy (relay forwarding). That strategy depends upon the fact that the address of the architected operand does not change during the life time of the instruction while it is in an active stations. For example, the architected addresses for register operands do not change for instructions. If the instruction takes as an input operand a register $r6$ for example, the address of the operand never changes for this particular instruction (it stays 6). This property is not generally true of memory operands. The difficulty with memory operands is that many memory related instructions determine the address of a memory operand value from an input register operand of the same instruction. Since we allow for instructions to execute and re-execute on entirely speculative operand values, the values of input register operands can be of essentially any value (including a wilding incorrect value) and thus the address of a memory operand can also change during the during that the instruction is in flight within the active station. This presents a problem for the correct enforcement of memory operands, and dependencies among them, in the entire program. If we examine the case of a memory store instruction, when it re-executes acquiring a new memory store value, the address of that memory store may also have changed ! We cannot simply forward that new memory operand (address and value) as with the relay forwarding strategy above. The reason is that we would not be superseding the previous memory operand that we forwarded previously that quite likely had a different architected address. Rather, we need some way to cancel the effect of any previously forwarded memory operands. This present forwarding strategy does just that.

In this strategy, memory operands that need to be forwarded employ a similar transaction as above for registers (described in the context of relay forwarding) but would instead have a transaction ID of *memory store* and would include the memory operand address and its value (along with the path and time-tag information). However, when an instruction either re-executes or its enabling predicate changes to being disabled, a different type of forwarding transaction is sent out. This new type of transaction is termed a *nullify transaction* and has the property of nullifying the effect of a previous store transaction to the same architected operand address. This transaction type consists of :

- a transaction ID of *memory nullify*
- the path ID on which this instruction is executing
- the time tag of the originating active station
- the memory operand address
- the value of the memory operand

When this transaction is snooped by subsequent ASes, for those ASes that have a memory operand as an input (that would be for instructions that load memory values in one way or another) a search is made for a match of an existing memory operand as usual but if a match is detected, the time-tag of that particular memory operand is set to a state such that any future *memory store* transaction, regardless of its time-tag value, will be accepted. Further, on reception of this *memory nullify* transaction, a request is sent backwards in program order for a memory

7

operand with the desired memory address. The transaction that represents a request for a memory operand would consist of :

- a transaction ID of *memory request*
- the path ID on which this instruction is executing
- the time tag of the originating active station
- the memory operand address

Of course, the memory address for the operand desired needs to be in the transaction but it is not as obvious why the originating AS's time tag is also included. In some interconnection fabrics, the time tag is included in backwarding requests to limit the scope of the travel of the transaction through the execution window. This same scope-limiting function is usually performed for forward going transactions as well. When the request is sent backwards in program order, previous ASes or the memory unit itself will eventually snoop the request and respond with another *memory store* transaction. As discussed, this forwarding strategy is very useful for memory operands but it can also be used for register operands with appropriate changes to the applicable transaction elements. Again, the inclusion of a time tag value is what allows for proper operand dependencies (whatever they may be for any particular forwarding strategy) in the committed program.

3.6.3 Predicate Forwarding

There are several ways in which instructions can be predicated in the microarchitecture. These predication mechanisms are not discussed in this paper but two such mechanisms are can be found in documents by Uht et al [26] and Morano [16]. For microarchitectures that predicate all program instructions within the microarchitecture itself (not visible at the ISA level of abstraction), predicate register values are essentially operands that need to be computed, evaluated, and forwarded much like register or memory operands. Each instruction computes its own enabling predicate by snooping for and snarfing predicate operands that are forwarded to it from previous instructions from the program-ordered past. Depending on the particular predication mechanism used, relay forwarding (described above) may be a suitable (if not good) choice for handling the forwarding of predicate operands. However, some predication mechanisms need additional transaction types (besides a base store transaction) to communicate. The predication mechanism described by Morano [16] requires three transaction to fully implement. That mechanism was employed for the data presented in a later section and the transactions for that mechanism are briefly described here.

This predication strategy requires two store-type transactions rather than just one. These two transactions are similar to other operand store transactions (like for register or memory operands) but one of these holds two values rather than just one. The first of these is the *region predicate store* transaction and consists of :

- a transaction ID of *region predicate store*
- the path ID on which this instruction is executing
- the time tag of the originating active station
- the region predicate value

This transaction is very analogous to a register or memory store but instead is used to forward a single bit value (the current *region predicate* for instructions following the AS that forwarded the transaction). A region predicate is a single bit that determines the execution status (enabled or disabled) for instruction that lie beyond the not-taken output path of a conditional branch. This particular transaction could be forwarded by either a conditional branch or by a an instruction that was not a control-flow-change instruction. In the case of a non-control-flow-change instruction, the only predicate value that makes sense is the same as its own enabling predicate and so only one value need be forwarded.

In the case of a conditional branch instruction, there are two possible output predicates that can be in view. One is for the not-taken output path from the branch. The other is for the taken output path. In order to forward both values for those instructions in program-ordered future, the other store transaction type (mentioned previously) us used. This transaction consists of :

- a transaction ID of *branch target predicate store*
- the path ID on which this instruction is executing

TT	instr	t1	t2	t3	t4
0	r3 := 1		r3=1		
1	r4 := r3 + 1			r3=1 r4=2	
2	r3 := 2	r3=2			
3	r5 := r3 + 2		r3=2 r5=4		

Figure 2: *Example Instruction Execution.* The time tags for sequential program instructions are on the left. Real time is show advancing along the top. For each real time interval, input operands are shown above any output operands.

- the time tag of the originating active station
- the branch target instruction address
- the region predicate value
- the branch target predicate value

This is identical to the previous *region predicate store* transaction but also includes the instruction address for the target of the conditional branch (the *taken* address) and the single bit predicate governing the execution status for instructions following the target of the conditional branch in program-ordered future.

Finally, for the predication mechanism employed in the present work, a third transaction is used to invalidate a previously forwarded branch target predicate. This transaction is a *branch target invalidation* and consists of :

- a transaction ID of *branch target invalidation*
- the path ID on which this instruction is executing
- the time tag of the originating active station
- the branch target instruction address
- the time tag of the branch target predicate to be invalidated

This is similar to other such invalidation transactions in that when it is snooped by ASes in the program-ordered future, a search is made for some state (in this case some predicate register state) that matches the given transaction criteria. The inclusion of the second time tag in this transaction allows for certain efficiencies that are particular to the predication mechanism described.

For predicate forwarding, as we have seen for register and memory forwarding, time tags play the vital role in identifying and preserving the ordering of all operands. In many ways, all operands (whether they be registers, memory, or execution predicates) require the use of time tags to determine the relative ordering of events in a microarchitecture that otherwise lets all instructions execute and re-execute wildly out of order in real time with respect to each other.

3.7   Execution Examples

A simple execution example using time tags is shown in Figure 2.

4   A Proposed Microarchitecture Using Time Tags

In this section we present a proposed microarchitecture that uses time tags as its basic dependency enforcing mechanism. The microarchitecture is ISA independent (can be applied to any ISA) and offer most of the features that are likely desirable in future machines. Among these are control-flow speculation and execution along with execution reuse of control-independent instructions beyond the joins of conditional branches, prediction of both
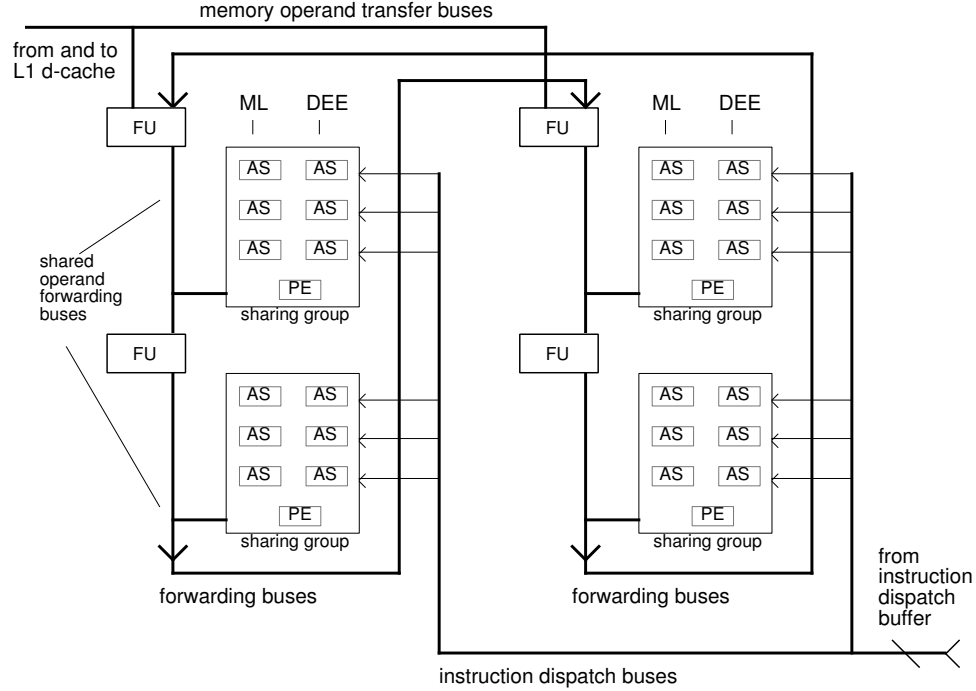
Figure 3: *Execution Window.* Shown are four sharing groups each with six active stations and one shared processing element. There are six rows of ASes in each column of the entire arrangement.

register and memory values, and multipath execution. Further, the proposed microarchitecture has a structure that allows for its size (numbers of various units) to scale without adverse implementation feasibility or significant performance loss.

4.1   Hardware Description and Operation

The most significant feature of this microarchitecture is that it makes use of an instruction issue slot that retains its decided instruction until it is ready to be retired (either committed or squashed) from execution. This is similar to what Lipasti and Shen employed in a microarchitecture proposed by them [15]. Our use of the issue slot is also similar to the old idea of reservation stations [24] since they are positioned in silicon in close proximity to execution resources. However, unlike reservation stations and more like an issue window, several of our instruction issue slots can share the same execution resources. We call our adaptation of the issue slot an *Active Station* (AS) where the idea of *active* illustrates that instructions dispatched to them will re-execute as necessary until retired.

Our Active Stations are laid out in silicon in a two dimensional grid whereby sequentially dispatched instructions (in program order) are assigned to sequential ASes down a column of the two dimensional grid of ASes. Time tags are thus assigned to the ASes in the same order that instructions were dispatched to them. As expected, these time tags have values starting at zero and incrementing up to one minus the total number of ASes in the machine.

Associated with each group of ASes are execution resources. We term each such resource a *Processing Element* (PE). A group of active stations along with their execution resource is termed a *Sharing Group* (SG). We also term the two dimensional grid of ASes, along with their interspersed execution units (PEs), the *Execution Window*. This arrangement is shown in Figure 3. There are four SGs shown, each with six ASes and one PE. The total height of each column consists of six ASes and there are two columns of SGs and four columns of ASes. In this microarchitecture, instructions are dispatched to an entire column of ASes simultaneously. At a maximum, new instructions can be dispatched to a column every cycle. Logically, the columns form a ring with the lower valued time tags on the left and the higher valued time tags on the right. As a column of instructions retires, that column

10

receives newly dispatched instructions that take on time tag values that will become the highest values in the machine. However all time tags are decremented when a new set of instructions is dispatched to a column. Thus the logical renaming of columns (and their constituent ASes) is effected. In addition to the ASes receiving time tags for tracking program order, all operands in the microarchitecture are also tagged with time tags. The usually comparisons (discussed previously) of time tags and operands architected addresses provides the basic dependency enforcing mechanism. Instructions generally re-execute when they snarf new input operands but they are also allowed to use value prediction for input operands not yet acquired through snooping.

This microarchitecture also takes advantage of the fact that instructions are dispatched to the ASes in program order down each column and continuing at the top of the next column to its right. Due to this physical ordering of the ASes, operand results only need to be forwarded to those ASes that lie below it, in its own column, and to those ASes in columns logically to the right. This arrangement provides for the opportunity to have a scalable interconnect fabric connecting the SGs to each other since there is no real need to provide connectivity from any single AS to prior ASes in the machine. The interconnection fabric for this microarchitecture is also segmented and both RFUs and MFUs are employed. Although not shown in Figure 3, three separate sets of forwarding buses were used for each of the three types of operands (register, memory, and predicate).

5 Simulation Results and Discussion

Using an execution-driven simulator, we ran SpecInt-2000 and SpecInt-95 programs on our time-tagged microarchitecture. Our goal here was to evaluate the instruction per clock (IPC) that was possible using the microarchitecture. Ten benchmark programs were used in all. Seven programs are from the SpecInt-2000 suite and three are from the SpecInt-95 suite. These programs were chosen in order to get a variety of execution behaviors. The microarchitecture simulated supports the MIPS-1 ISA (big endian), with some MIPS-2 instructions also supported in order to accommodate code residing in the SGI system libraries which use them. All programs were compiled using the vendor SGI compiler on the SGI IRIX 6.4 operating system. Programs were compiled with standard optimization (-O) for primarily the MIPS-1 ISA (-mips1). The first 600 million instructions of all programs were executed. Data were only gathered after the execution of the first 100 million instructions (a total of 500 million). The default features of the machine simulated are given in Table 1. The data in Table 2 contains IPC results for a range of machine sizes. All machines simulated also contained two forwarding buses for register operands, two forwarding buses for predicates, and one forwarding bus for memory operands. All forwarding buses incur a bus transfer delay of 1 clock. Further, each forwarding unit encountered (different for different machine sizes) has a minimum latency of 1 clock. All simulated machines, regardless of size, have forwarding buses that span eight sharing groups. This illustrates how the bus length can be constant with respect to the machine size thus allowing for physical scalability of the machine. Machine sizes are basically characterized by Each of the machine configurations in Table 2 consists of three numbers that give the rows of sharing groups, the number of AS rows per sharing group, and the total number of columns respectively. The number of sharing group rows times the number of ASs per sharing group is the total number of AS rows in a configuration. The product of all three numbers gives the total number of ASes in the machine and therefor the number of instructions that may be in flight simultaneously (having possibly speculatively executed already). Generally, as the number of ASs increases, the resulting IPCs also increase, but there are diminishing returns. A significant IPC gain is achieved, for example, when increasing the size of the machine from the 8-4-8 geometry (256 ASes) to the 8-8-8 geometry (512 ASs). This is an increase in IPC of approximately 17.8% on the harmonic mean across all benchmarks. However, doubling the number of ASs again (the number of instructions in flight in the e-window simultaneously) to 1024 with the 16-8-8 machine geometry, the harmonic mean IPC only increases by about 3.4%. Even an alternative geometry (the 8-16-8 geometry) also having 1024 ASs gives a harmonic mean IPC of 4.96, an increase of approximately 5.6% over the 8-8-8 geometry. Comparing the 8-16-8 geometry with the 16-8-8 geometry (both having 1024 ASs), the results from each are fairly similar with the 8-16-8 geometry edging out the 16-8-8 when all benchmarks are considered. However, for some benchmarks (like GCC and VORTEX) the IPC are lower with the 8-16-8 geometry. This suggests that there is little benefit in just increasing the number of AS rows per sharing group after some point (16 in the present) case since contention for the single common processing logic starts to limit performance. With the current state of this microarchitecture research, the best tradeoff of resources and IPC results would appear to be with the machine

Table 1: *General machine characteristics.* These machine parameters are used for all simulations as the default except where one of these parameters may be varied.

L0 cache size	32 words
L1 I/D cache access latency	1 clock
L1 I/D cache size	64 KBytes
L1 I/D block size	32 bytes
L1 I/D organization	2-way set associative
L2 cache access latency	10 clocks
L2 cache size	2 MBytes
L2 block size	32 bytes
L2 organization	direct mapped
main memory access latency	100 clocks
forwarding unit minimum latency (all)	1 clock
forwarding-bus latency (all)	1 clock
number of register forwarding buses	2
number of predicate forwarding buses	2
number of memory buses	1
branch predictor	PAg
	1024 PBHT entries
	4096 GPHT entries
	saturating 2-bit counter

Table 2: *Benchmark IPC results for various machine sizes.* Different machine sizes are characterized by their geometries consisting of the three-number entries along the top of the table: SG rows per column, AS rows per SG, and SG columns.

config	8-4-8	8-4-12	12-4-8	8-8-8	8-12-8	8-16-8	16-8-8	32-8-8
bzip2	4.1	4.1	4.3	4.9	5.2	5.5	4.9	5.1
compress	4.1	4.3	4.3	4.7	4.8	4.8	4.8	4.9
crafty	3.7	3.8	3.8	4.1	4.2	4.1	4.1	4.4
gcc	3.1	3.1	3.1	3.3	3.3	3.3	3.4	3.6
go	4.1	4.6	4.7	5.1	5.5	5.4	5.3	6.0
gzip	4.8	5.0	5.2	6.1	6.5	6.4	6.4	6.5
ijpeg	6.2	7.2	7.8	9.5	12.3	13.1	12.0	12.3
mcf	3.5	3.5	3.6	4.3	4.7	5.3	4.5	5.0
parser	3.5	3.7	3.5	3.9	4.0	3.9	3.9	4.4
vortex	4.2	4.4	4.7	4.8	4.9	4.6	4.9	5.1
H-MEAN	3.8	4.2	4.3	4.7	4.9	5.06	4.9	5.2

Table 3: *Benchmark IPC results for various machine sizes when assuming 100% branch prediction accuracy guiding i-fetch.* Different machine sizes are characterized by their geometries consisting of the three-number entries along the top of the table: SG rows per column, AS rows per SG, and SG columns.

config	8-4-8	8-8-8	16-8-8	32-8-8
bzip2	9.0	12.4	15.5	16.3
compress	8.1	12.2	16.2	17.6
crafty	6.1	9.3	14.7	21.7
gcc	7.1	12.3	17.3	21.2
go	9.0	12.3	17.3	21.2
gzip	7.6	10.0	11.2	11.5
ijpeg	9.2	14.4	18.7	20.2
mcf	6.1	8.2	10.1	12.0
parser	7.3	10.9	14.6	17.0
vortex	6.4	10.1	15.2	21.7
H-MEAN	7.3	10.8	13.3	16.9

geometry of 8-8-8. This consists of 512 ASs and the associated bus interconnects between them. Without the use of time tags as an operand dependency enforcing mechanism, it is not clear how this arrangement could otherwise be implemented in current or near-term process technologies due to interconnection and contention problems that would arise.

A better look at the potential of this microarchitecture is presented by the IPC data in Table 3. Some of the configurations presented previously, in Table 2, are repeated but this time they were executed with a relaxed assumption about the branch prediction accuracy guiding the fetching and dispatching of instructions into the e-window. In this data set, 100% branch prediction accuracy was assumed in an attempt to see what might be possible with this general microarchitectural approach as research continues to make possible i-fetch improvements. In this IPC data, results continue to improve up through the 32-8-8 machine geometry. Any realization of a machine with this number of speculative instructions in flight (2048 in this case) is clearly not feasible given current operand enforcement methods like a reorder buffer. Contention for access to it would be too great given current and foreseeable silicon limitations.

6   Summary

We have presented a method for tracking and enforcing program dependencies through the use of time tags. Register, memory, and predicate dependencies can all be managed with time tags. We have also described several microarchitectural features are either enabled or become more feasible through the use of time tags for dependency enforcement. We have also presented a proposed microarchitecture that used time tags throughout and which allows for control-flow speculation and reuse of control-flow independent instructions, data speculation, and multipath execution. The proposed microarchitecture is also physically scalable through the use of a segmented operand forwarding fabric. The use of time tags allows for a degree of out-of-order execution that is not easily enforceable using any other mechanism due to either routing congestion or access contention problems. We also presented results that indicates that this general approach appears to be quite promising as compared with the existing more conventional program dependency enforcement mechanisms. Some work on much larger machine configurations has already suggested that achieving IPC numbers in the 10s on general integer sequentially-oriented program codes is possible.

References

[1] Anderson D., Sparacio F. and Tomasulo F. The IBM/360 Model 91: Machine Philosophy and Instruction Handling. *IBM Journal*, 11(1):8–24, Jan. 1967.

[2] V. T. Cher C. Skipper: A microarchitecture for exploiting control-flow independence. In *Proceedings of the 34nd International Symposium on Microarchitecture*, New York, NY, Nov 2001. ACM Press.

[3] Cleary J.G, Pearson M.W and Kinawi H. The Architecture of an Optimistic CPU: The Warp Engine. In *Proc. of the HICSS*, pages 163–172, Jan. 1995.

[4] P. B. et al. Internal architecture of alpha 21164 microprocessor. In *COMPCON'95*, pages 79–87, Mar 1995.

[5] S. G. Franklin M. Arb: A hardware mechanism for dynamic memory disambiguation. *IEEE Transactions on Computers*, 45(5):552–571, May 1984.

[6] S. G. Franklin M. Register traffic analysis for streamlining inter-operation communication in fine-grained parallel processors. In *Proceedings of the 25th International Symposium on Microarchitecture*, pages 236–245, New York, NY, Dec 1992. ACM Press.

[7] Gonzalez J. and Gonzalez A. Limits on Instruction-Level Parallelism with Data Speculation. Technical Report UPC-DAC-1997-34, UPC, Barcelona Spain, 1997.

[8] Heil T.H., Smith J.E. Selective Dual Path Execution. Technical report, University of Wisconsin - Madison, Madison, WI, 1996.

[9] G. J. and G. A. Limits of instruction level parallelism with data speculation. In *Proceedings of the VECPAR Conference*, pages 585–598, 1998.

[10] W. D. Kessler R.E., McLellan E.J. The alpha 21264 microprocessor architecture. In *International Conference on Computer Design*, Oct 1998.

[11] Klauser A.S., Pathankar A., Grunwald D. Selective eager execution on the polypath architecture. In *Proceedings of the 25th International Symposium on Computer Architecture*, pages 250–259, New York, NY, Jun 1998. ACM Press.

[12] Lam M.S. and Wilson R.P. Limits of Control Flow on Parallelism. In *Proc. of ISCA-19*, pages 46–57. ACM, May 1992.

[13] M. H. Lipasti and J. P. Shen. Exceeding the dataflow limit via value prediction. In *International Symposium on Microarchitecture*, pages 226–237, 1996.

[14] M. H. Lipasti and J. P. Shen. The performance potential of value and dependence prediction. In *European Conference on Parallel Processing*, pages 1043–1052, 1997.

[15] Lipasti M.H and Shen J.P. Superspeculative Microarchitecture for Beyond AD 2000. *IEEE Computer*, 30(9), Sep 1997.

[16] Morano D.A. Execution-time Instruction Predication. Technical Report TR 032002-0100, Dept. of ECE, URI, Dec 2002.

[17] Palacharla S., Jouppi N.P. and Smith J.E. Complexity-effective superscalar processors. In *ISCA-24*, pages 206–218, 1997.

[18] S. K. Sankaranarayanan K. A scheme for selective squash and re-issue for single-sided branch hammocks. Technical Report CS-2001-14, University of Virginia, Charlottesville, VA, Jul 2001.

[19] S. K. Sankaranarayanan K. A scheme for selective squash and re-issue for single-sided branch hammocks. IEEE Press, Oct 2001.

[20] A. P. Skadron K. Hydrascalar: A multipath-capable simulator. *IEEE Technical Committee on Computer Architecture Newsletter*, 2001.

[21] S. G. Smith J.E. The microarchitecture of superscalar processors. In *Proceedings of the IEEE*, volume 83, pages 1609–1624, Dec 1995.

[22] B. S. Sohi G.S. and V. T.N. Multiscalar processors. In *Proceedings of the 22th International Symposium on Computer Architecture*, New York, NY, Jun 1995. ACM Press.

[23] Thornton J.E. Parallel Operation in the Control Data 6600. In *Proc. of the AFIPS*, volume 26, pages 33–40, 1964.

[24] Tomasulo R.M. An Efficient Algorithm for Exploiting Multiple Arithmetic Units. *IBM Journal of Research and Development*, 11(1):25–33, Jan 1967.

[25] Uht A. K. and Sindagi V. Disjoint Eager Execution: An Optimal Form of Speculative Execution. In *Proc. MICRO-28*, pages 313–325. ACM, Nov 1995.

[26] Uht A.K., Morano D.A., Khalafi A., de Alba M., Wenisch T., Ashouei M., and Kaeli D. IPC in the 10's via Resource Flow Computing with Levo. Technical Report TR 092001-001, Dept. of ECE, URI, Sept 2001.

[27] Wallace S., Calder B., Tullsen D.M. Threaded multiple path execution. In *Proceedings of the 25th International Symposium on Computer Architecture*, New York, NY, Jun 1998. ACM Press.

[28] Wang S.S.H, Uht A.K. Ideograph/ideogram: Framework/architecture for eager execution. In *Proceedings of the 23rd Symposium and Workshop on Microprogramming and Microarchitecture*, pages 125–134, New York, NY, Nov 1990. ACM Press.