



URI / NEU research collaboration

Levo ILP Machine

ILP Speed-Ups in the Tens !

student **David Morano**
advisors **Professor David Kaeli**
Professor Augustus Uht

NUCAR 01/02/22

Northeastern University Computer Architecture Research

Levo ILP Machine

DAM 1

outline



- **problem and introduction**
 - the teaser
- **background**
- **key concepts and high-level machine organization**
- **some major machine components**
 - sharing groups
 - active stations
 - register filters
 - memory filters
- **examples**
 - resource flow
 - branch predication
- **conclusions**
- **resources**

problem



- we want programs to run faster !
- we often want single threaded programs to run faster !
- our approach is Instruction Level Parallelism (ILP)
- we want existing ISAs and compiled programs to benefit (micro-architectural solution)
- we need to employ some or most of the following design principles :
 - we need to have a wide instruction window to expose the ILP
 - we need a wide issue machine to increase IPC
 - we need to manage and exploit minimal control dependencies
 - we need to perform control and data speculation
- we end up with the Levo machine ! (one way to get high ILP)
- ILP in the tens ! (this is our motto)



URI / NEU research collaboration

instruction level parallelism

- J. Fisher's paper (1981)
- dates from at least the CDC 6600 (Thornton 1964 !)
- Tomasulo came up with reservation stations around 1967 !

What are the major components of machine speed ?

$$\frac{\text{work}}{\text{time}} = \frac{\text{work}}{\text{instr}} * \frac{\text{instrs}}{\text{clock}} * \frac{\text{clocks}}{\text{time}}$$

speed ————

technology

micro-architecture

architecture

IPC = instrs / clock (key measurement of ILP)

tough program codes



easy program to parallelize

```
input a[] ;  
input b[] ;  
for (long time) {  
    c[i] = func(a[i]) ;  
    d[i] = func(b[i]) ;  
    i += 1 ;  
}  
output c[] , d[] ;
```

harder program to parallelize

```
input a ;  
for (long time) {  
    b = func(a) ;  
    c = func(b) ;  
    a = func(c) ;  
}  
output a ;
```

- + of course, we want to try and parallelize the hard programs !!
- + there are already several ways to deal with easier programs

the big teaser !



**Lamb and Wilson, "Limits of Control Flow on Parallelism" ISCA 19th, 1992
(Stanford)**

**using the Spec benchmarks: awk, ccom, eqntott, espresso, gcc, irsom, latex,
matrix300, spice2g6, tomcatv they evaluated the limits on the amount of
possible parallelism in the programs**

some upper-bound results :

mean speed-up	execution constraints
1.0	instructions execute one after the other
2.14	an instruction cannot execute until its immediately preceeding branch is resolved
6.80	instructions execute when their immediately preceeding mispredicted brach is resolved
39.62	instructions execute when their mispredicted control dependent branches (all of them) resolve
158.26	all branches are perfectly predicted and all instructions are allowed to speculatively execute ahead (Oracle)

teaser -- but !



**in order to get the big gains, we must have a machine that allows
(according to Lamb and Wilson) :**

- the elimination of all but true data dependencies**
- speculative execution beyond just a speculatively executed branch**
- independent regions of code must be allowed to execute speculatively and simultaneously (minimal control dependencies)**
- use of some sort of predicated instruction execution (guarding)**
- execution of multiple paths simultaneously !**

background

- **1972 Garold S. Tjaden** - detection of concurrency and reduced control dependencies
- **1986 Uht** - extracting hardware concurrency
- **1991 Popescu, et al** - time tags for squashing on a branch misprediction
- **1991 Uht** - theory of minimal control dependencies
- **1995 Cleary, Pearson, Kinawi** - Warp Engine, used "time-stamps" for state roll-back of speculatively executed "events"
- **1995 Sohi, Breach, Vijaykumar** - Multiscalar Processor, concurrently executing pieces of a single path execution, uses compilation assistance communicated through the machine ISA
- **1995 Uht, Sindagi** - disjoint eager execution (DEE)
- **1997 Uht, Sindagi, Somanathan** - branch effect reduction techniques
- **1997 Uht** - verification of ILP speedups for DEE (TR)
- **1997 Uht** - high performance memory for ILP (TR)
- **1997 Uht** - Levo high-ILP computer (TR)

levo history

- **Levo I**
 - used dependency matrices for data and control to determine when to execute instructions
- **Levo II**
 - introduced resource flow computing (free-for-all speculative execution)
 - introduced "hidden" hardware predication of all instructions (predicate forwarding)
 - but was not scalable
- **Levo III**
 - introduced scalability through the use of $O(\text{constant})$ bus spans
 - used centralized register files for architected registers
 - used an ARB type memory strategy
 - no memory value forwarding yet
- **Levo IV (current machine model)**
 - distributed persistent store of architected registers
 - backward buses for requesting register values
 - memory value forwarding and backwarding
 - memory write store queue for committed writes

building a dream machine

- **elimination of all name dependencies through renaming**
 - register & memory
- **wild speculative execution without regard for any dependencies**
 - control or data
- **all instructions are predicated at instruction load-time in hardware**
- **scalable machine to support a large execution window**
 - buses broken up by forward or filter units
- **real-estate consumming resources (like floating point) can be shared**
- **execution of multiple paths simultaneously with DEE (equal probability wave-front execution)**

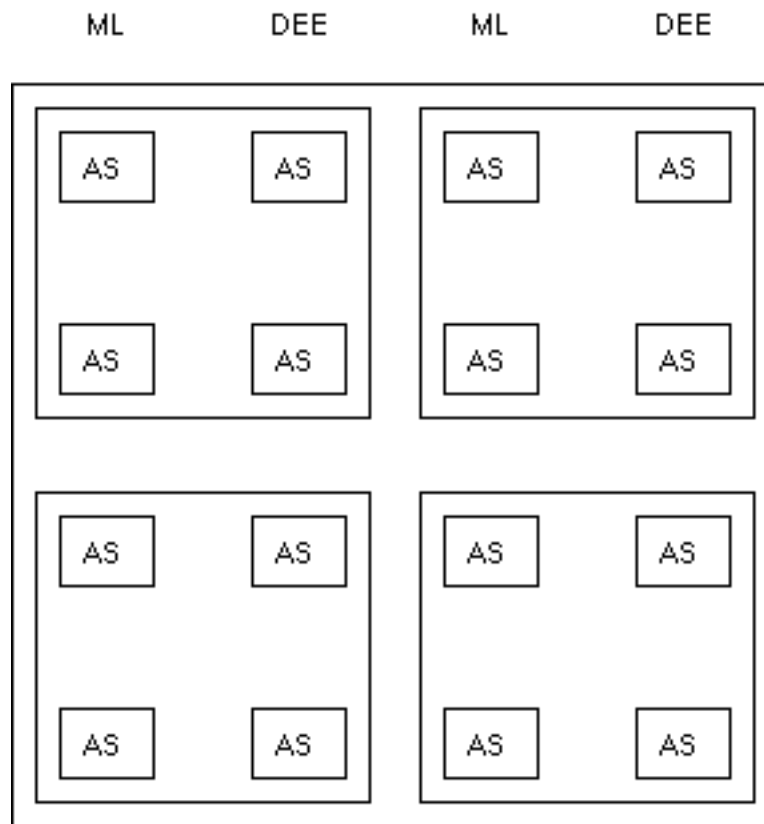
key Levo concepts

- **disjoint eager execution (old)**
- **resource flow computing (new)**
 - execution proceeds primarily on resource availability rather than data or control dependencies !
- **hardware branch predication (new)**
- **wide instruction window and issue (old)**
 - "A 21st century microprocessor may well [issue] up to dozens of instructions [per cycle, peak] ..." -- Dave Patterson, 1995
- **active station concept (new -- extension of a reservation station)**
 - instructions re-execute as necessary until retired
- **execution (resource) sharing groups (new and old)**
- **result forwarding buses -- extended Common Data Buses (Tomasulo)**
- **high-bandwidth interleaved memory interface (old)**
 - redundant load/store elimination
- **scalability (new for this scale of an ILP machine)**
- **memory and register filters units (scalability and bandwidth conservation)**

high-level block diagram



8 ML Active Stations
8 DEE Active Stations
2 ML columns
2 DEE columns
4 Sharing Groups



instruction execution window

I-fetch and
branch prediction

branch tracking
buffer

station load buffer

write store queue

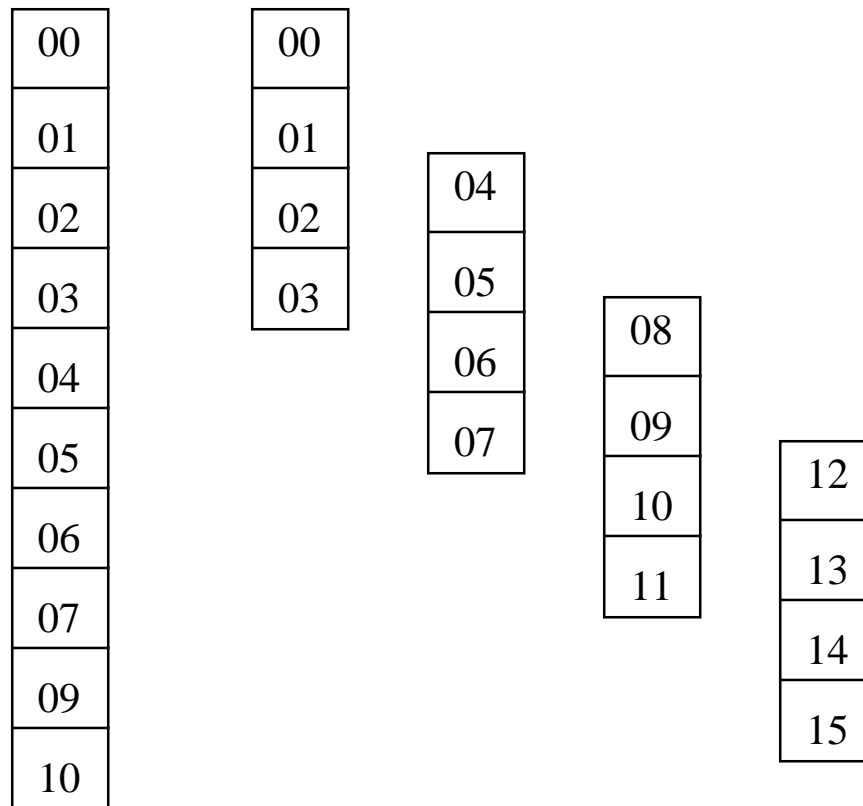
memory caches

main memory

- **execution flow**

- fetch
- load
- issue
- dispatch
- execute
- re-execute !
- retire

folded execution window



00	04	08	12
01	05	09	13
02	06	10	14
03	07	11	15

- **logical view of instructions**
- **columns form groups of instructions that are retired together**
- **peak simultaneous load is a whole column**
- **peak simultaneous execution is the number of sharing groups in the machine**

logical and physical groupings



logical
organization

mainline (ML) path

00	04	08	12
01	05	09	13
02	06	10	14
03	07	11	15

DEE
path 1

00	04
01	05
02	06
03	07

DEE
path 2

00	04
01	05
02	06
03	07

physical
organization

ML D1

00	00
01	01
02	02
03	03

ML D1

04	04
05	05
06	06
07	07

ML D2

08	00
09	01
10	02
11	03

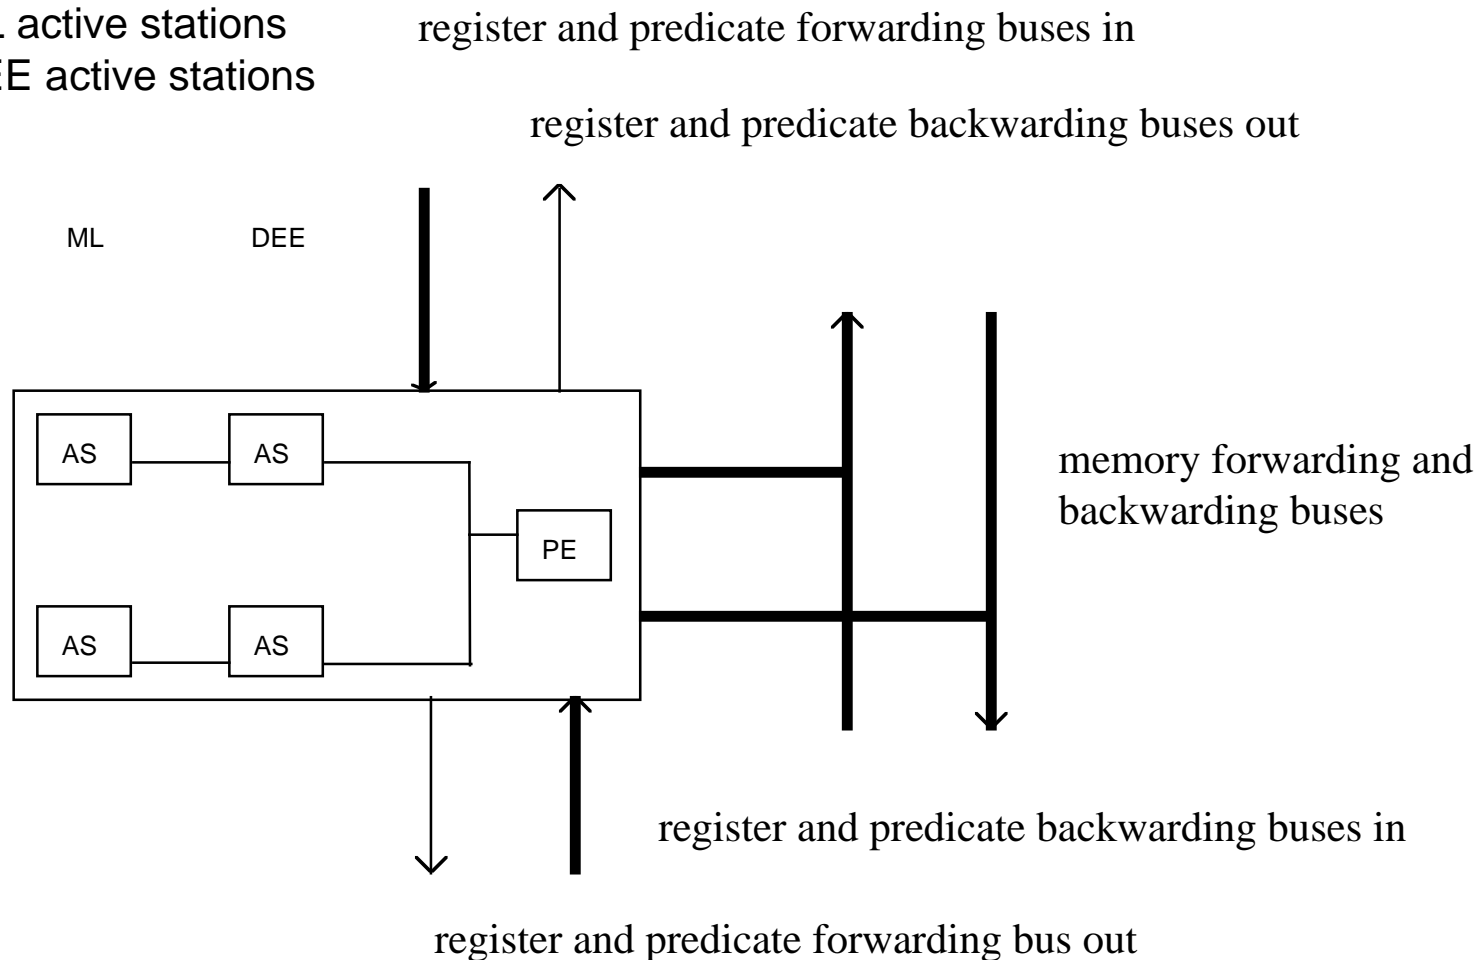
ML D2

12	04
13	05
14	06
15	07

sharing group



- 1 sharing group
- 2 ML active stations
- 2 DEE active stations

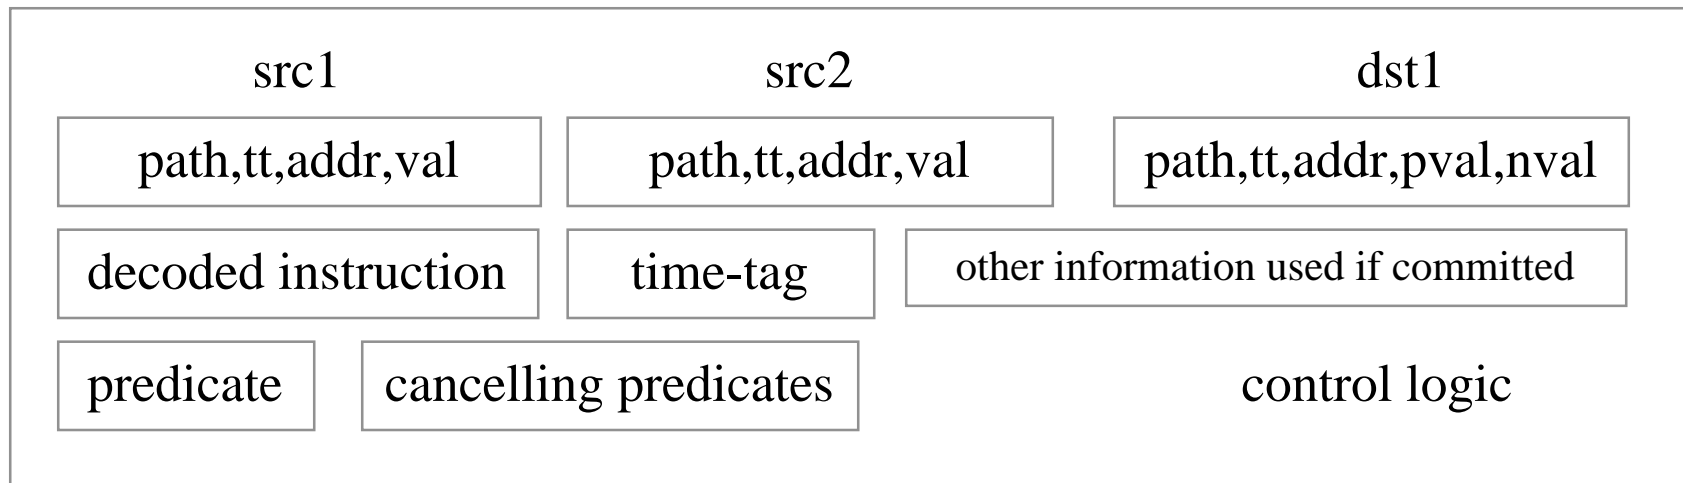


active station

- sort of an extension to the reservation station (Tomasulo)
- serves as a place to wait for operands and control flow changes
- holds an instruction until it is ready to retire
- snoops all result forwarding buses coming to it for new operands
- makes requests for values on backwarding buses (register & memory)
- snoops for control flow (prediction) changes by watching predicates representing the branch domains that the instruction is in
- re-executes !! an instruction when one or more of its operands change
- can squash its own execution with a "relay" operation (restores its output data as if it had not executed)
- can rebroadcast its output on a switch from a DEE path to the ML path


active station

- similar to Tomasulo's reservation station
- implements dynamic register renaming



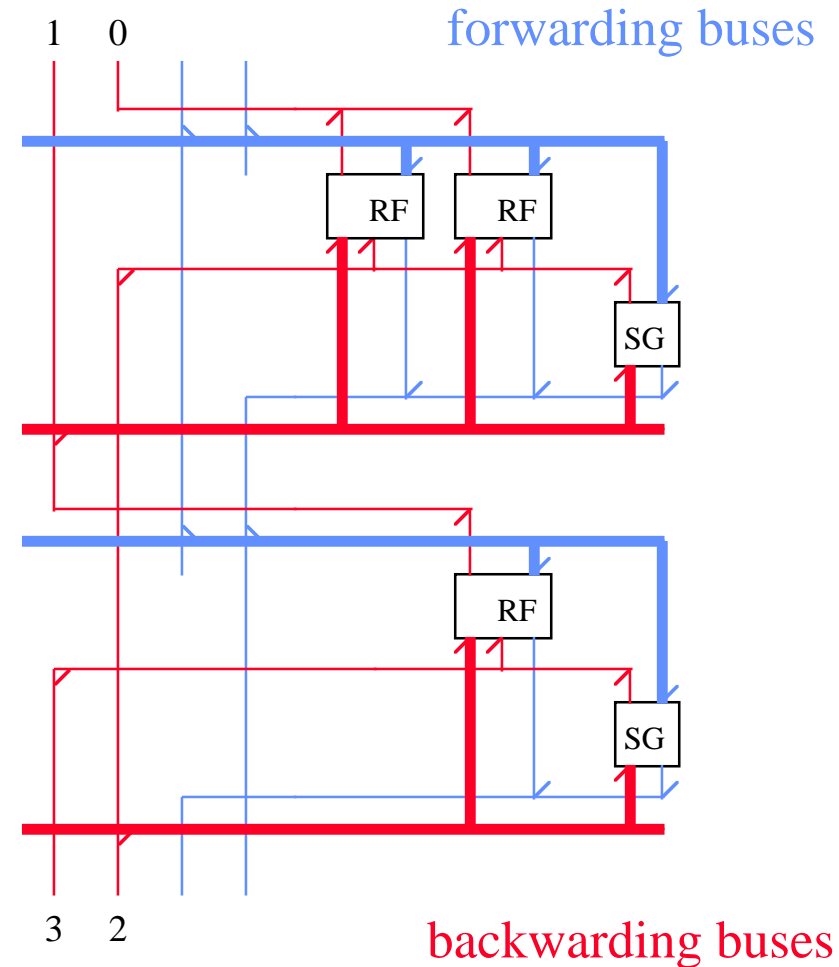

- register names take the form of:
 - path : time-tag : addr
 - for example "1 : 27 : r6"

register filter unit



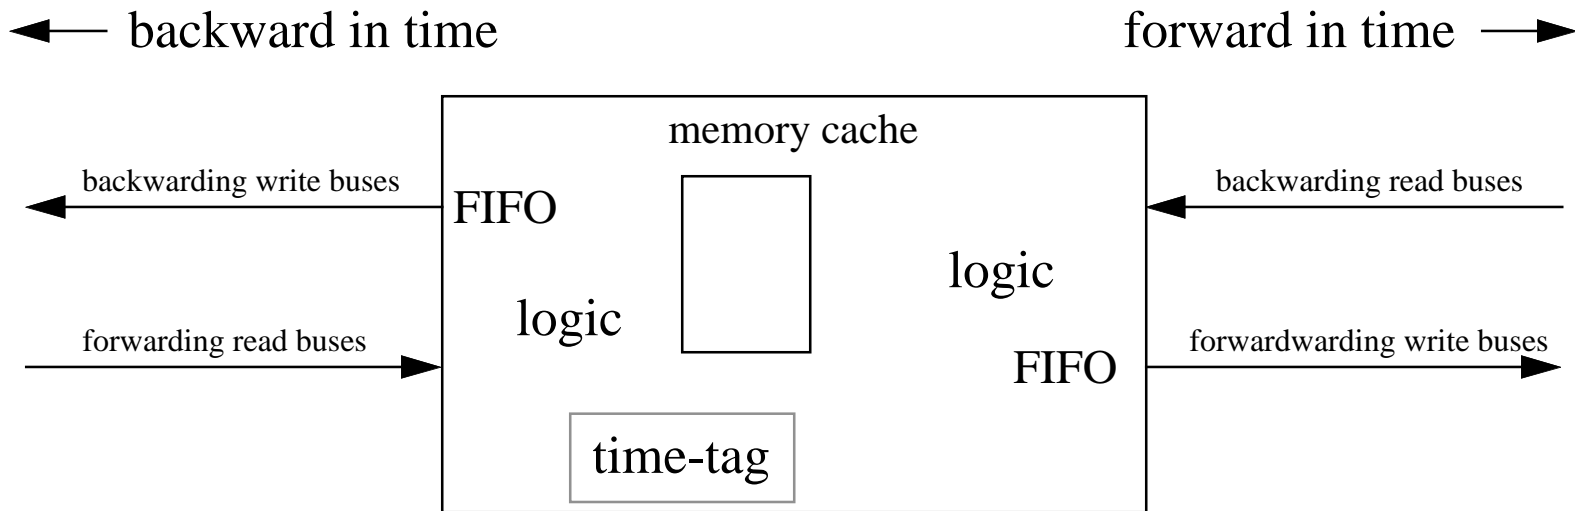
- all buses are register transaction buses
- consolidates update transactions on input
- updates are "forwarded" to the output bus request logic immediately when possible
- updates are aged in the file store so that the oldest ones do not get starved out from being forwarded (oldest goes first) ; backward requests are aged similarly

interconnections



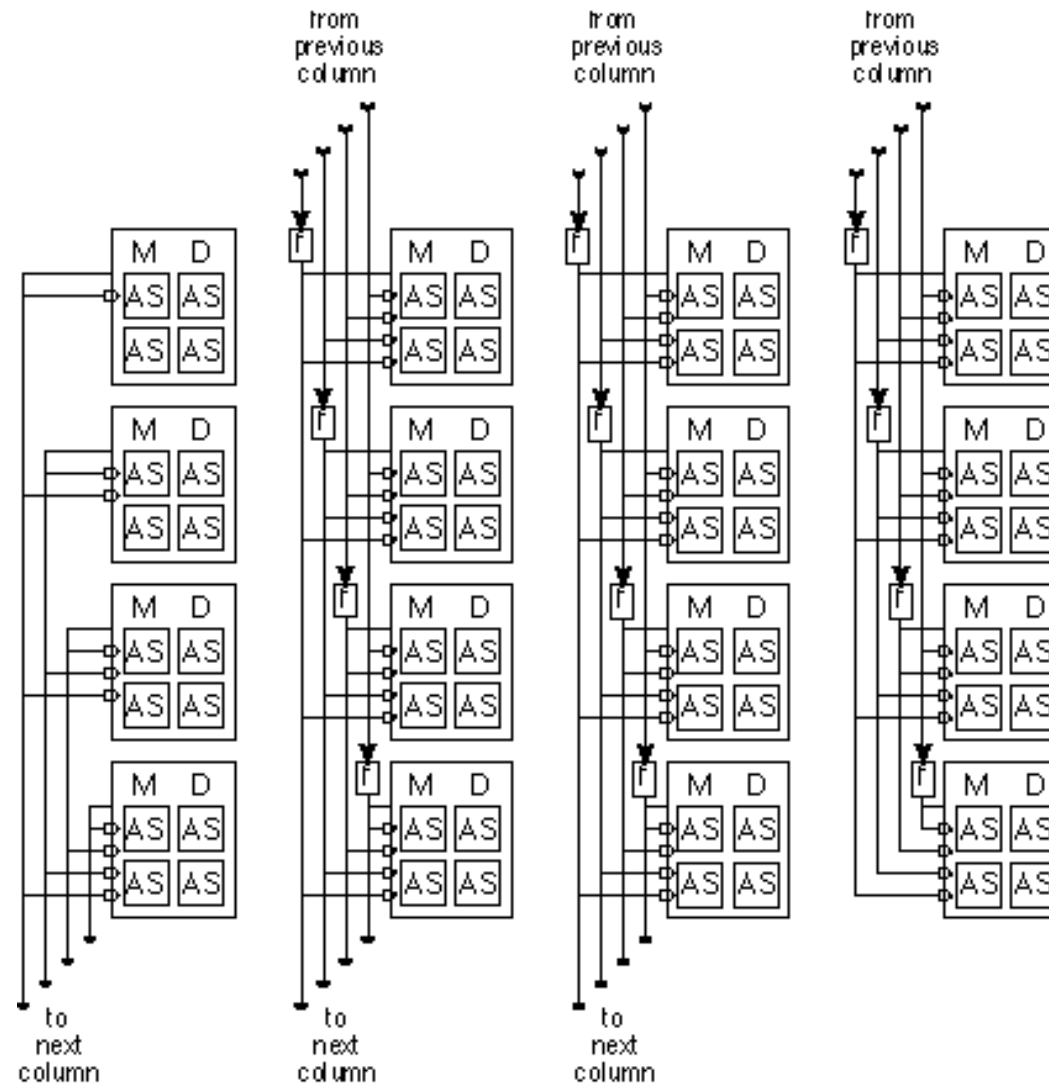
forwarding & backwarding span of 2

memory filter unit



- all buses are memory buses (number of which set according to interleave factor)
- consolidates update transactions on input
- updates are "forwarded" to the output bus request logic immediately when possible
- current policy is to queue outgoing requests or responses in FIFOs until the buses are granted for use

sharing group forwarding buses



resource flow example



- **X** -- a ML execution
- **R** -- relay operation
- **B** -- broadcast only operation
- **D** -- disjoint path execution

			time					
instruction	00	r9 <- r0 op r1	X					
	10	r3 <- r5 op r6	X					
	20	r3 <- r7 op r8	X					
	30	r2 <- r3 op r8	X	X				

- all execute immediately
- instruction 30 re-executes when it gets an updated version of r3
- instruction 30 ignores the output from 10 instead preferring that from 20

resource flow example



- **X** -- a ML execution
- **R** -- relay operation
- **B** -- broadcast only operation
- **D** -- disjoint path execution

			time					
instruction	00	r9 <- r0 op r1	X					
	10	r3 <- r5 op r6	X					
	20	r3 <- r7 op r8			X			
	30	r2 <- r3 op r8		X		X		

- 00 and 10 execute immediately, instructions 20 and 30 delayed for some reason
- instruction 30 executes when it gets r3 from 10
- instruction 20 eventually gets to execute
- instruction 30 re-executes with the output from 20

resource flow example



- **X** -- a ML execution
- **R** -- relay operation
- **B** -- broadcast only operation
- **D** -- disjoint path execution

			time					
instruction	00	r3 <- r5 op r6		X				
	10	r1 <- r3 op r1	X		X			
	20	r3 <- r5 op r6	X					

- 10 and 20 execute immediately, instruction 00 delayed
- instruction 10 does not re-execute due to its own broadcast of register r1 or because of the broadcast of r3 from instruction 20
- instruction 00 eventually executes
- instruction 10 re-executes because of forwarded value from 00

resource flow example



- **X** -- a ML execution
- **R** -- relay operation
- **B** -- broadcast only operation
- **D** -- disjoint path execution

		time						
		0	1	2	3	4	5	6
instruction	00		X				X	
	10	X		X				X
	20				X			
	30		X	X		X		

- 10 gets to execute first for whatever reason
- 00 and 30 then get to execute
- instructions 10 and 30 re-execute because of forwarded value from 00
- 20 finally gets to execute
- 30 will re-execute because of the update from 20
- 00 re-executes for whatever reason causing 10 to re-execute but none later

resource flow example



- forwarding span delays
- X -- a ML execution
- R -- relay operation
- B -- broadcast only operation
- D -- disjoint path execution

			time						
			0	1	2	3	4	5	6
instruction	000	r2 <- r1 op r0	X						
	040	r3 <- r2 op r0	X		X				
	080	r4 <- r2 op r0	X			X			
	120	r5 <- r3 op r0	X				X		

- assume a forwarding span of 32 instructions
- all instructions get to execute immediately
- r2 output from 000 is latched and only reaches 040 after a clock delay due to span distance
- the same r2 from 000 reaches instruction 080 after a forwarding span delay (1 clock) also
- 120 executes due to updated value (r3) from 040 also after a forwarding span delay

branch predication example



- minimal control dependency
- instruction relaying
- X -- a ML execution
- R -- relay operation
- B -- broadcast only operation
- D -- disjoint path execution

			time						
			0	1	2	3	4	5	6
instruction	00	r2 <- r1 op r0	X						
	10	b_op r2, 030	X	X					
	20	r3 <- r1 op r0	X		R				
	30	r4 <- r1 op r0	X						

- the branch is initially predicted to be NOT taken !
- all instructions get to execute immediately
- instruction 10 re-executes due to new value from 00
- the branch changes its prediction due to being re-executed !
- instruction 20 (in the branch domain) forwards the original value for r3
- instruction 30 does not have to re-execute even though the branch prediction changed !

branch predication example



- minimal control dependency
- instruction relaying
- time-tag comparisons
- X -- a ML execution
- R -- relay operation
- B -- broadcast only operation
- D -- disjoint path execution

			time						
			0	1	2	3	4	5	6
instruction	00	r2 <- r0 op r1			X				
	10	b_op r2, 030				X			
	20	r2 <- r3 op r0	X				R		
	30	r4 <- r2 op r0		X				X	

- the branch is initially predicted to be NOT taken !
- instruction 20 executes because it is in the predicted path
- instruction 30 executes due to new value from 20
- 00 finally executes causing the branch to re-execute bit not 30 ! (later in time)
- the branch is now predicted taken causing instruction 20 to relay
- the relayed value from 20 causing 30 to re-execute

branch predication example



- DEE execution example
- minimal control dependency
- instruction relaying
- time-tag comparisons
- X -- a ML execution
- R -- relay operation
- B -- broadcast only operation
- D -- disjoint path execution

			time						
			0	1	2	3	4	5	6
instruction	00	$r2 \leftarrow r0 \text{ op } r1$	X				X		
	10	$b_op \ r2, 030$	X	X				X	
	20	$r3 \leftarrow r0 \text{ op } r1$	R		D				B
	30	$r4 \leftarrow r0 \text{ op } r1$	X						

- the branch is predicted to be taken at the start of this example
- all instructions are executed immediately because they are in the main-line path, 20 is a relay
- 10 re-executes due to updated value from 00
- instruction 20 executes in a DEE path some time later
- 00 re-executes causing 10 to re-execute resolving the branch to not-taken and the DEE path to switch and become the new ML path, this causes instruction 20 to perform a broadcast

conclusions

- we realize disjoint eager execution (DEE)
 - we realize resource flow execution
 - hardware branch predication
 - wide instruction window and issue
 - active stations are the key to instruction re-execution as necessary
 - result forwarding buses (extended CDBs) provide for data updates and scalability
 - the machine is constant and linearly scalable with instruction window, issue, and execution resources !
-
- the future of ILP is
 - the future of speculative execution is
 - the future of computing is



URI / NEU research collaboration

Levo resources

- **Professor A. Uht (University of Rhode Island)**
 - <http://www.ele.uri.edu/~uht/>
 - has much of the relatively recent published and unpublished work
- **Levo information page**
 - <http://ovel.ele.uri.edu:8080/>
 - currently the main source for Levo information
- **LevoVis**
 - <http://ovel.ele.uri.edu:8080/levovis.html>
 - visual display simulator for the Levo I machine
- **this talk**
 - <http://www.ele.uri.edu/~morano/levo/>



URI / NEU research collaboration